

Sliced Designs for Multi-platform Online Experiments

Soheil Sadeghi

Microsoft, sosadegh@microsoft.com

Peter Chien

Department of Statistics at University of Wisconsin-Madison, peterq@stat.wisc.edu

Neeraj Arora

Wisconsin School of Business at University of Wisconsin-Madison, neeraj.arora@wisc.edu

December 15, 2017

Multivariate testing is a popular method to improve websites, mobile apps, and email campaigns. A unique aspect of testing in the online space is that it needs to be conducted across multiple platforms such as a desktop and a smartphone. The existing experimental design literature does not offer precise guidance for such a multi-platform context. In this paper we introduce a multi-platform design framework that allows us to measure the effect of the design factors for each platform and the interaction effect of the design factors with platforms. Substantively, the resulting designs are of great importance for testing digital campaigns across platforms. We illustrate this in an empirical email application to maximize engagement for a digital magazine. We introduce a novel “sliced effect hierarchy principle” and develop design criteria to generate factorial designs for multi-platform experiments. To help construct such designs, we prove a theorem that connects the proposed designs to the well-known minimum aberration designs. We find that experimental versions made for one platform should be similar to other platforms. From the standpoint of real world application, such homogeneous sub-designs are cheaper to implement. To assist practitioners, we provide an algorithm to construct the designs that we propose.

Key words: A/B testing; design of experiments; blocking; digital marketing; email testing; web experiments

1. Introduction

A highly desirable characteristic of marketing activities online is that they lend themselves to rapid and extensive testing. Online retailers routinely improve the layout of their website to maximize profitability. Websites intended for the purpose of educating visitors evaluate user engagement on an ongoing basis and attempt to improve metrics that include page views and time spent per page. The simplest form of online experimentation is A/B testing. It has become quite popular because tools such as Google Analytics make it easy to implement. A more informative form of online testing is multivariate testing because it allows one to assess the individual effect of multiple factors at the same time. The primary focus of this paper is on constructing new designs for multivariate testing for online experiments.

Online testing is a popular and rapidly growing method in the digital world for applications that include website and email optimization. A report by Forrester evaluates a variety of suppliers that include Adobe, Maxymiser and Optimizely which are well known in the online testing area (Stanhope 2013). For website optimization, consider, for example, the homepage of Liberty Mutual Insurance shown in Figure 1. A typical A/B testing may begin by creating two versions of the website: version A and version B. To implement the test, the website traffic is divided into two disjoint sets. Engagement metrics such as the number of pages navigated determine the winner of the test.

A more informative form of testing involves multiple attributes of the website. In a very simple example, consider the following four website factors, each at two levels (present versus absent): the login tab on top right, the ‘Resources tab’ below it, ‘find a local agent’ tab in the middle, and the ‘Auto Savings and Discounts’ visual at bottom. To find the best combination of these four attributes, one can create sixteen versions of the website for a full factorial design. For an engagement metric such as page views, this multivariate test can help uncover the incremental contribution of each factor and as a result help improve the website layout. As the number of attributes increase, a fraction of all combinations could be used in order to perform the multivariate test (Wu and Hamada 2011). This is often necessary to ensure feasibility and to address cost constraints.

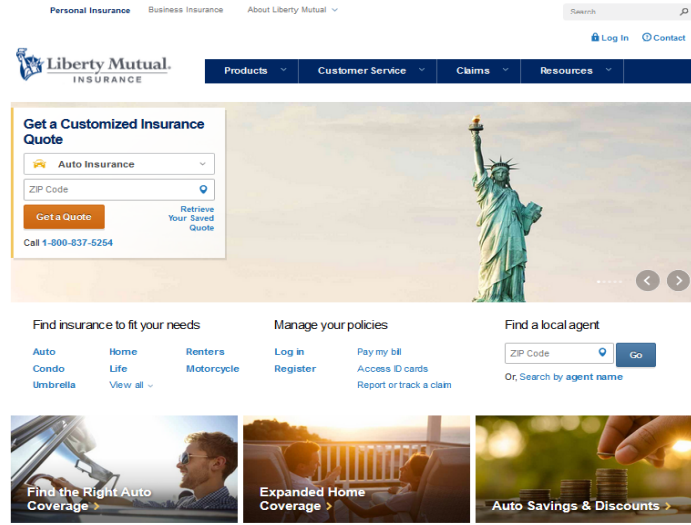


Figure 1 Liberty Mutual Insurance Homepage.

Multivariate experiments in the online space present a new design challenge: experiment needs to be conducted across multiple platforms that include desktops, laptops, tablets, and smartphones. Such experiments are important because a different set of attribute combinations may be optimal for each platform. For example, presence of multiple images may be best for a desktop and in contrast a list of links may be more effective for a smartphone. In this paper, we introduce a new design, called the *sliced factorial design*, for such multi-platform, multivariate online experiments. We develop optimality criteria to construct such designs. This complete design we propose can be partitioned into sub-designs or slices such that each slice is used for one platform.

We have chosen to use the term “slice” because this is how it has been used in the design literature historically. Qian and Wu (2009) pioneered the idea of *sliced designs* in the context of computer experiments. Qian (2012) showed how the idea could be effectively combined with Latin hypercube sampling. A sliced Latin hypercube design in Qian (2012) can be divided into slices of smaller Latin hypercube designs. Such a design is attractive for building emulators for computer models and performing numerical integration. Similarly, sliced designs constructed by Xu et al. (2011) and Qian (2012) are intended for running simulation experiments and building Gaussian process prediction models. Follow-up work on sliced designs for computer experiments include Yang et al. 2013, Huang et al. 2014, Yin et al. 2014, and Ba et al. (2015).

In contrast with extant work, the sliced factorial design proposed by us is intended for performing multi-platform experiments in the digital space. The sliced designs for computer experiments are focused on low-dimensional projections whereas our sliced designs explore *slicing* in a factorial design set up.

We begin the paper with a review of factorial designs at two levels. Factorial designs are built upon three fundamental and widely used principles: effect hierarchy, effect sparsity, and effect heredity (Wu and Hamada 2011). In our design framework we ensure that each slice or sub-design follows these fundamental principles. In order to account for the unique multi-platform context, we generalize the effect hierarchy principle to the *sliced effect hierarchy* principle. Our main idea is that each sub design allows accurate assessment of factors that are significant for the corresponding platform and the complete design allows us to identify the interaction between the design factors and the platforms. We impose a reasonable ordinal structure on the effects the experiment is trying to uncover in the cross-platform context. Based on the sliced effect hierarchy principle, we develop extended design criteria called *sliced resolution* and *sliced aberration*. These criteria are then used to generate the sliced factorial designs to perform a multi-platform, multivariate test.

Minimum aberration designs are popular among practitioners and tables to construct them are readily available in software and textbooks (Wu and Hamada 2011). In an effort to build upon the rich results of the aberration literature, we prove a theorem that connects sliced factorial designs to minimum aberration designs. Acting as the bridge between what we already know, this theorem helps construct sliced factorial designs that we propose. Using the extended design criteria and this theorem, we develop an algorithm to construct sliced factorial designs. This algorithm is for the case where all the combinations of design factors are feasible for all the platforms and works for both symmetric and asymmetric designs. We also provide a procedure for the more general situations where some combinations of design factors may be infeasible for some platforms. This incorporates a typical design constraint wherein not all factors combinations may be feasible for all platforms. An example of such a design constraint on a smartphone would be the inability to

include two image based factors, a brand logo and the picture of a child, on one screen. The same combination may work quite well on a desktop.

Based on our new design criteria, we find that it is desirable to have the sliced factorial design be divided into homogeneous slices: experimental versions made for a slice should be as similar as possible to the ones for other slices. In terms of resources, the multivariate design we propose requires a smaller number of versions compared to other designs. We illustrate that although slices are used to uncover factor effects within each platform, the sliced factorial design can be used to compare the results across different platforms.

Finally, we illustrate our novel design framework in the context of an empirical email optimization application intended to maximize engagement for a digital magazine. Algorithm 1 was used to generate a sliced factorial design for this study. The results reveal interesting insights about how different factors were effective for different platforms. Based on the results for this multi-platform experiment, the expected gain in page views for the two platforms was 16% and 7%.

The remainder of the paper is organized as follows. In Section 2 we lay a foundation for our work by reviewing key design concepts upon which we build our framework. In Section 3 we formulate the multi-platform experiments and explain the sliced hierarchy principle. The following section introduces two design criteria that are useful to construct the new designs we propose and presents designs' construction. In Section 5, we present an empirical application as an illustration for our design framework. Finally, we offer our conclusions in the last section of the paper.

2. Review of Factorial Designs at Two Levels

Factorial designs are often used in studies where the interest is to model the effects of more than one factor simultaneously. These designs are well studied in statistics. Wu and Hamada (2011) and Box et al. (1978) are excellent references on this subject. To serve as a basis for further development, in this section we provide a brief review of factorial designs and the relevant concepts. For ease of presentation, we begin with the case that includes k two-level factors, denoted by $1, \dots, k$. For each factor, its linear effect on a response variable, denoted by y , is investigated.

A *full factorial design* requires 2^k runs for the k factors. Denote \mathbf{y} as the vector of responses of length 2^k , and \mathbf{X} as the $2^k \times 2^k$ model matrix for which the 2^k columns consist of the column of ones \mathbf{x}_0 , k columns of ± 1 values $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_k$ representing the design vectors, $\binom{k}{2}$ columns of two-way interactions $\mathbf{x}_{12}, \mathbf{x}_{13}, \dots$ (each a product of two design columns), $\binom{k}{3}$ columns of three-way interactions $\mathbf{x}_{123}, \mathbf{x}_{124}, \dots$ (each a product of three design columns), ..., all the way to the column of k -way interaction $\mathbf{x}_{12\dots k}$ (the product of all k design columns).

The regression model can be written as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta}, \quad (1)$$

which does not include the error term as this is a fully saturated model, with the same number of coefficients as number of observations. The regression coefficient $\boldsymbol{\beta}$ consists of elements β_0 (the intercept), β_1, \dots, β_k (main effects), $\beta_{12}, \beta_{13}, \dots$ (two-way interaction effects), ..., $\beta_{12\dots k}$ (the k -way interaction effect). Obviously, not all effects need to be included in the model. For instance, a model of main effects only includes k main effects and a noise component that reflects all the interaction effects. The least square estimate of the regression coefficient $\boldsymbol{\beta}$ is given by

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}, \quad (2)$$

where \mathbf{X}' is the transpose of the matrix \mathbf{X} , and $(\mathbf{X}'\mathbf{X})^{-1}$ is the inverse of the matrix $\mathbf{X}'\mathbf{X}$. For a 2^k full factorial design, all diagonal elements of $\mathbf{X}'\mathbf{X}$ are equal to 2^k , and off-diagonal elements are zero. Therefore, the estimate of an element of $\boldsymbol{\beta}$, e.g. $\hat{\beta}_{ij\dots l}$, is given by

$$\hat{\beta}_{ij\dots l} = \frac{1}{2^k} \sum_{t=1}^{2^k} c_t y_t = \frac{1}{2} \left(\frac{1}{2^{k-1}} \sum_{t=1}^{2^k} c_t y_t \right), \quad (3)$$

where the weights c_t 's are the ± 1 elements of the corresponding vector column $\mathbf{x}_{ij\dots l}$. For a 2^k full factorial design, half of the elements of each column $\mathbf{x}_{ij\dots l}$ (excluding the column of ones \mathbf{x}_0) are equal to $+1$, and the rest half of elements are equal to -1 . The model matrix (excluding the column of ones) of a full 2^3 factorial design, provided in the first seven columns of Table 1, is an example. Hence, the estimate $\hat{\beta}_{ij\dots l}$ can be written as

$$\hat{\beta}_{ij\dots l} = \frac{1}{2} \left(\frac{1}{2^{k-1}} \sum_{t=1}^{2^k} c_t y_t \right) = \frac{1}{2} \left(\bar{y}(\mathbf{x}_{ij\dots l}+) - \bar{y}(\mathbf{x}_{ij\dots l}-) \right), \quad (4)$$

where $\bar{y}(\mathbf{x}_{ij\dots l}+)$ is the average of the y_t values at (+) level of column $\mathbf{x}_{ij\dots l}$, and $\bar{y}(\mathbf{x}_{ij\dots l}-)$ is the average of the y_t values at (-) level of column $\mathbf{x}_{ij\dots l}$. In the factorial design literature, usually $\bar{y}(\mathbf{x}_{ij\dots l}+) - \bar{y}(\mathbf{x}_{ij\dots l}-)$ is reported as the estimate, denoted by **ij ... k** effect, and not the regression coefficient estimate $\hat{\beta}_{ij\dots l}$, i.e.,

$$\mathbf{ij} \dots \mathbf{k} \text{ effect} = \bar{y}(\mathbf{x}_{ij\dots l}+) - \bar{y}(\mathbf{x}_{ij\dots l}-). \quad (5)$$

For example, in Table 1, main effect **1** is estimated by $\frac{1}{4}(y_5 + y_6 + y_7 + y_8) - \frac{1}{4}(y_1 + y_2 + y_3 + y_4)$, and interaction effect **23** is estimated by $\frac{1}{4}(y_1 + y_4 + y_5 + y_8) - \frac{1}{4}(y_2 + y_3 + y_6 + y_7)$.

Table 1 Model Matrix and Data for 2^3 Design

1	2	3	12	13	23	123	Data
\mathbf{x}_1	\mathbf{x}_2	\mathbf{x}_3	\mathbf{x}_{12}	\mathbf{x}_{13}	\mathbf{x}_{23}	\mathbf{x}_{123}	\mathbf{y}
—	—	—	+	+	+	—	y_1
—	—	+	+	—	—	+	y_2
—	+	—	—	+	—	+	y_3
—	+	+	—	—	+	—	y_4
+	—	—	—	—	+	+	y_5
+	—	+	—	+	—	—	y_6
+	+	—	+	—	—	—	y_7
+	+	+	+	+	+	+	y_8

Full factorial designs are rarely used in practice for large k . Instead, a fraction of a full factorial design, called a *fractional factorial design* denoted by 2^{k-p} , is often used. In general, 2^{k-p} denotes a $(\frac{1}{2})^p$ fraction of a 2^k factorial design. The optimal fraction can be selected according to resolution and aberration based criteria (Wu and Hamada 2011). A design can be replicated more than once and a design with a single replicate is called an unreplicated design.

2.1. Properties and Principles of Factorial Designs

Two key properties of factorial designs are *balance* and *orthogonality*. A design is balanced if each factor level appears in the same number of experimental runs. Two factors are called orthogonal if all their different combinations appear in the same number of experimental runs. A design for which all pair of its factors are orthogonal is called an orthogonal design. As one would expect, a full factorial design is balanced and orthogonal.

After defining the properties, an important question is to understand the relative importance and relationship between the effects. Below we list three fundamental principles that serve as the foundation for factorial designs (Wu and Hamada 2011, p. 172).

1. **Effect Hierarchy:** This principle indicates that lower-order effects are more likely to be important than higher-order effects. For example, main effects are more likely to be important than interaction effects. Further, effects of the same order are equally likely to be important.
2. **Effect Sparsity:** It indicates that number of relatively important effects in a factorial design is small.
3. **Effect Heredity:** In order for an interaction to be significant, at least one of its parent main effects should be significant.

These fundamental principles guide the construction and analysis of factorial designs. In most situations, the effects of lower order interactions are believed to be more important than the higher order effects. For this reason, a fractional factorial design can be generated by confounding the effects of higher order interactions with the lower order ones. Lower order effects therefore can be estimated by assuming that higher order interactions are negligible.

3. Formulation of Design for Multi-Platform Experiments

Multivariate testing is a popular method in the digital world for applications that include website and email optimization. However, multivariate experiments in the online space present a new design challenge: experiments need to be conducted across multiple platforms that include desktops, laptops, tablets, and smartphones. Such multivariate experiments are important because a different set of attribute combinations may be optimal for different platforms.

A unique aspect of the multi-platform context is that in addition to the design factors one also has to pay attention to the platform variable. The purpose of any multivariate experiment is help select the optimal level for each design factor. This is in contrast to the platform variable, where the goal is optimize for each level of the platform variable. That is, unlike a design factor, for the platform variable we are not trying to select an optimal level. This presents a design challenge that

this paper is trying to solve. As it relates to the analysis of the results from such a multi-platform experiment, the proposed design allows us to first analyze the data for each platform. Next, the same data across the platforms allows one to test for interactions between the platform variable and the factors.

Our basic idea is that the statistical design in the multi-platform context should permit (i) an assessment of factor effects for each platform and (ii) the interaction effect between each factor and the platform variable. The design we propose accomplishes both these goals. Later in the paper we explain how existing experimental designs fail to adequately incorporate the platform variable in the multi-platform context. As far as we are aware, this is the first systematic statistical work in this direction.

Earlier we noted the importance of effect hierarchy when designing an experiment. A fundamental challenge for a multi-platform design is how to think about the platform variable in the effect hierarchy. We incorporate the following order of effects in the multi-platform experiment we design. To begin, the response variable of interest (e.g. page views) is expected to vary by platform. Accurate assessment of difference in page views by platform is an important first step before we assess which factors are significant and for which platform. It is necessary to control for between-platform difference in page views to be able to accurately assess factor effects. This is easily seen in a context where certain factors are only significant for one platform that gets lower web traffic. In such a context, failure to accurately assess between-platform differences in page views could very easily mask factor effects for the smaller platform.

After accounting for the between-platform difference effect, the effects associated with the design factors are the central focus of the multi-platform experiment. As only some of the design factors may be significant, accurate assessment of the magnitude of the significant design factors is important. Finally, the effects of the design factors likely vary by platform. The design needs to accurately assess such platform-by-design factor interactions. Among all the effects that an experiment can uncover, our focus as it relates to the design, is on the following effects and in this order:

the between-platform difference effect, the design factor effects and the interaction between the two. We formalize the above design ideas next.

To facilitate development of the ideas we propose, we formally define a multi-platform experiment as follows.

DEFINITION 1 (Multi-platform Experiment). *Consider a multi-platform experiment for studying k two-level design factors, denoted by $1, \dots, k$, on s platforms P_1, \dots, P_s . The complete design set d of the experiment consists of s sub designs, d_1, \dots, d_s , with d_j associated with P_j . To quantify the difference among the platforms, let S denote a categorical factor, called the slice factor, with s levels. The j th level of S is associated with P_j .*

Figure 2 visually displays the design set d of the experiment in Definition 1.




Figure 2 Design Set d of the Experiment in Definition 1.

To construct suitable designs for the experiment in Definition 1, we introduce two guiding properties:

PROPERTY 1. For $j = 1, \dots, s$, the sub design d_j should achieve desirable estimation capacity¹ for the design factors on platform P_j .

PROPERTY 2. Combined together, the complete design d should achieve desirable estimation capacity for the slice factor S and the two-way interactions between S and the design factors.

¹ Please refer to Cheng et al. (1999) for a formal definition of estimation capacity.

In a standard factorial design, the effect hierarchy principle implies that a lower-order effect is more likely to be important than a higher-order effect, and that all effects of the same order are equally likely to be important. Therefore, to construct the sub design d_j , it is desirable to focus on main effects of the design factors. This follows Property 1. Further, the focus of the complete design d should be on the main effect of the slice factor S and two-way interactions between S and the design factors. This follows Property 2.

The set of effects in a multi-platform experiment are expected to follow an order that we formalize next. We formally state the *sliced effect hierarchy principle* for d to incorporate the slice factor S in the design construction process. To do so, we divide the effects of d into two disjoint sets.

For the design d in Definition 1, let E_I be the set of all effects with words that exclude the slice factor S and E_S be the set of all effects with words that include the slice factor S . Using this notation, we define the sliced effect hierarchy principle as follows:

Sliced Effect Hierarchy Principle

- (i) For E_I or E_S , the lower-order effects are more likely to be important than the higher-order effects.
- (ii) For E_I or E_S , effects of the same order are equally likely to be important.
- (iii) Any effect in the set E_S is likely to be more important than an effect in E_I that is of the same order.
- (iv) Any effect in the set E_S is likely to be less important than an effect in E_I that is of a lower order.

3.1. Illustration

Next we use a simple example to illustrate the sliced effect hierarchy principle and how it differs from the effect hierarchy principle defined in Section 2.

EXAMPLE 1. For the experiment in Definition 1, let $k = 3$ and $s = 2$ for two platforms P_1 and P_2 . The slice factor S with two levels, $-$ and $+$, represents P_1 and P_2 , respectively. Sub designs d_1

and d_2 are factorial designs with three factors such that each includes seven factorial effects that are ranked in Table 2 following the effect hierarchy principle.

Table 2 Effect Hierarchy Principle for Each Sub Design of the Experiment in Example 1

Rank	Effects
(i)	1, 2, 3
(ii)	12, 13, 23
(iii)	123

The complete design d is a factorial design with three design factors and the slice factor S , and includes fifteen factorial effects(see Table 3). The two sets E_I and E_S are $\{1, 2, 3, 12, 13, 23, 123\}$ and $\{S, 1S, 2S, 3S, 12S, 13S, 23S, 123S\}$, respectively. Following the sliced effect hierarchy principle defined above, Table 3 ranks all fifteen factorial effects.

Table 3 Sliced Effect Hierarchy Principle for the Complete Design of the Experiment in Example 1

Rank	E_S	E_I
(i)	S	
(ii)		1, 2, 3
(iii)	1S, 2S, 3S	
(iv)		12, 13, 23
(v)	12S, 13S, 23S	
(vi)		123
(vii)	123S	

Example 1 illustrates that the sliced effect hierarchy principle assumes a certain order of effects for the design. The main effect **S** is at the top of the hierarchy to account for the possibility of platform differences for the dependent variable. This is an important difference to account because factor main effects may be significant for low traffic platform but not for the high traffic platform. Recent industry reports (e.g. Adobe Digital Index) emphasize the appropriateness of such a hierarchy by documenting the substantial gap in page views by platform type. One such report (Chaffey 2016) notes that a majority of site visits across a wide variety of industries are from desktops thus indicating that the between-platform difference could be large and significant. Our empirical application reported later further illustrates this point.

Next in the sliced effect hierarchy are the main effects of the design factors (**1, 2, 3**), the primary focus of interest. The slice factor precedes the design factors because the effects of design factors on the dependent variable is conditioned on the level of the platform variable. This is followed by the interactions between the design factors and the slice factor (**1S, 2S, 3S**). The interactions between the design factors and the slice factor are important because the effects of design factors on the dependent variable usually varies in different levels of the platform variable. This is followed by two-way interactions between design factors (**12, 13, 23**), and the remaining higher order interactions (**12S, 13S, 23S; 123; 123S**).

As we will show next, the sliced effect hierarchy principle guides the design that is best suited for a multi-platform experiment.

3.2. Sliced Effect Hierarchy and Existing Design Methods

The sliced effect hierarchy is a new design principle and as a result, existing methods cannot be used to construct designs that conform with it. We illustrate this point by discussing how some design methods that appear to be reasonable choices fail to work in the multi-platform contexts. First, consider a random splitting approach that would begin by generating a factorial design for d and then randomly split it into sub designs for d_j 's. Although simple to implement, such an approach cannot guarantee that each d_j follows Property 1. Second, another method could be to independently generate all sub designs d_j 's and then put them together to form d . Unfortunately, such an approach cannot guarantee that d has a desirable structure according to Property 2. Third, blocking is a standard method to form blocks of homogeneous units in a factorial design. While this method works well for agriculture and engineering applications where treatment-blocking interaction is negligible (Wu and Hamada 2011), it is ill-suited for our new design problem. If one uses the slice factor S as a block factor to construct a blocked factorial design d with blocks d_1, \dots, d_s , then S would be confounded with the higher-order interaction effects of the design factors. This assumes that the slice factor S has a negligible interaction with the design factors. This assumption contradicts Property 2 and is counter to the primary goal of how design factor effects may interact with platforms.

In view of the drawbacks of the aforementioned methods, in the next section we explain how to construct new designs for the experiment in Definition 1. Our basic idea is that each sub design d_j should follow the effect hierarchy principle and the complete design d should follow the sliced effect hierarchy principle defined above. We construct the designs by extending resolution and aberration based criteria and call them *sliced factorial designs*.

4. Construction of Sliced Factorial Designs

In this section, we propose a method for constructing the sliced factorial designs. To do so requires extending the maximum resolution and minimum aberration criteria to accommodate the slice factor. For ease of presentation, we consider the experiment in Definition 1 with two platforms. Assume that the complete design d is a 2^{k+1-p} fractional factorial design that consists of two sub designs, d_1 and d_2 , each of which is a 2^{k-p} fractional factorial design. The maximum resolution and minimum aberration criteria, as defined in Section 2 for standard fractional factorial designs, is used to judge the statistical quality of each sub design. As noted earlier, the resolution based criterion maximizes the *length* of the shortest word in a defining relation of a factorial design and the aberration based criterion minimizes the *number* of shortest words. For the complete design d , we need to modify these optimality criteria to accommodate the sliced effect hierarchy principle. To accomplish this, we first provide some useful notation and definitions.

DEFINITION 2. For a 2^{k+1-p} complete design d in Definition 1 with $s = 2$,

1. The sliced defining relation is obtained by multiplying the defining relation of d by the slice factor S .
2. A sliced word is any word in the sliced defining relation except for the slice factor S .
3. The sliced wordlength pattern is the vector

$$SW = (3^{B_3}, \dots, (k+1)^{B_{k+1}}) \quad (6)$$

where B_j is the number of sliced words of length j and k is the number of design factors.

We now revisit Example 1 to illustrate Definition 2.

EXAMPLE 1 (CONTINUED). Building upon this example that we set up earlier, assume that the complete design is a 2^{4-1} fractional factorial design that consists of two sub designs, each of which is a 2^{3-1} fractional factorial design. The design $d^{(1)}: \mathbf{I} = \mathbf{123S}$ is a minimum aberration design with four factors that can be used as the complete design of this multi-platform experiment. As a result of $d^{(1)}$, the sub designs on P_1 and P_2 are $d_1^{(1)}: \mathbf{I} = -\mathbf{123}$ and $d_2^{(1)}: \mathbf{I} = \mathbf{123}$, respectively. Both $d_1^{(1)}$ and $d_2^{(1)}$ are minimum aberration designs with three factors. For the complete design $d^{(1)}$, the sliced defining relation is $\mathbf{S} = \mathbf{123}$ which is obtained by multiplying $\mathbf{I} = \mathbf{123S}$ by the slice factor S . The design $d^{(1)}$ has one sliced word $\mathbf{123}$, which is of length three. The sliced wordlength pattern is then $SW(d^{(1)}) = (3^1)$.

We now extend resolution and aberration to accommodate the slice factor. We define *sliced resolution* as follows:

Sliced Resolution: *The sliced resolution of a 2^{k+1-p} complete design d in Definition 1 with $s = 2$ is defined to be the smallest sr such that $B_{sr} \geq 1$, i.e., the length of the shortest sliced word in the sliced defining relation.*

Following the sliced effect hierarchy principle, one is interested in maximizing the sliced resolution. Maximizing the sliced resolution ensures that the slice factor S and its interaction with the design factors can be best estimated. Further, for situations where the complete design d cannot be judged by its sliced resolution alone, we define *sliced minimum aberration* as follows:

Sliced Minimum Aberration Designs: *Suppose that, for the experiment in Definition 1 with $s = 2$, two 2^{k+1-p} complete designs $d^{(1)}$ and $d^{(2)}$ are to be compared. Let sr be the smallest integer such that $B_{sr}(d^{(1)}) \neq B_{sr}(d^{(2)})$. Design $d^{(1)}$ is said to have less sliced aberration if $B_{sr}(d^{(1)}) < B_{sr}(d^{(2)})$. If there is no design with less sliced aberration than $d^{(1)}$, then $d^{(1)}$ is called a sliced minimum aberration design.*

For a given k and p , a sliced minimum aberration design always exists. Built upon the sliced effect hierarchy principle, the sliced minimum aberration criterion allows any two complete designs to be ranked.

Sliced resolution and sliced aberration based criteria are defined with the goal of extending resolution and aberration to accommodate the sliced effect hierarchy principle. Developed on the effect hierarchy principle, the resolution based criterion maximizes the length of the shortest word and the aberration based criterion minimizes the number of shortest words in a defining relation of a factorial design. To extend these concepts to accommodate the sliced effect hierarchy principle, the sliced resolution based criterion maximizes the length of the shortest sliced word and the sliced aberration based criterion minimizes the number of shortest sliced words in a sliced defining relation of a factorial design. Working with sliced words in a sliced defining relation, instead of words in a defining relation, ensures that the slice factor S and its interaction with the design factors can be best estimated using the complete design (Property 2).

Having defined a suitable design criterion for a multi-platform context, we are now ready to construct sliced minimum aberration designs. Theorem 1 below establishes a connection between the sliced minimum aberration criterion and the well-established minimum aberration criterion. Our intention here is to guide the construction of the sliced minimum aberration designs using readily available minimum aberration designs of fewer number of factors.

THEOREM 1. *A 2^{k+1-p} sliced minimum aberration design corresponds to a sliced defining relation in which all sliced words contain S .*

Proof. A sliced defining relation with all sliced words containing S means the slice factor S is not involved in any word of the defining relation for the sliced minimum aberration design. Therefore, no generator uses the slice factor S to generate the sliced minimum aberration design. For fixed k and p , it is sufficient to show that any sliced defining relation with at least a sliced word excluding S is inferior to a sliced defining relation that includes S in all its sliced words. Any sliced factorial design with a sliced word excluding S has S involved in the defining relation of the sliced factorial design, and hence S is used in the generators to generate the sliced factorial design. We prove for the case where one generator uses S . The proof can be easily generalized to the case where more than one generator uses S . Consider a 2^{k+1-p} sliced factorial design d that has $p - 1$ generators

not involving S and one generator g involving S . It is sufficient to show that a design with all generators excluding S is better according to the sliced aberration criterion. Form a new design d_{new} by removing S from g . Call the new generator g_{new} . As S only appears in g , the product of g_{new} with other generators will result in a word excluding S . Comparing the sliced defining relation of d with d_{new} 's shows that the length of all the sliced words formed by g_{new} has increased by one, and the lengths of all other sliced words remain the same. Therefore, d_{new} is better according to the sliced aberration criterion.

□

We now revisit Example 1 to illustrate Theorem 1.

EXAMPLE 1 (CONTINUED). Let us revisit the design $d^{(1)}$ that we discussed earlier. We note that $d^{(1)}$ with the sliced defining relation $\mathbf{S} = \mathbf{123}$ is not a sliced minimum aberration design because the sliced word $\mathbf{123}$ does not contain S . The only generator of $d^{(1)}$, as defined in Section 2, is $\mathbf{123S}$ which contains S . Following the proof of Theorem 1, removing S from this generator results in a sliced defining relation of $\mathbf{S} = \mathbf{123S}$ that is superior to the sliced defining relation of $d^{(1)}$.

To find a sliced minimum aberration design, it is sufficient to search among possible complete designs for which the sliced defining relations contain S in all their sliced words. As all sliced words contain S , removing S from the sliced defining relation of such designs result in a defining relation with all words excluding S . Therefore, minimizing the number of shortest sliced words in the sliced defining relation of a 2^{k+1-p} complete design d is equivalent to minimizing the number of shortest words in the defining relation of a 2^{k-p} fractional design consisting of design factors only. In other words, The sliced defining relation of a 2^{k+1-p} sliced minimum aberration design can be generated by multiplying the slice factor S to the defining relation of a 2^{k-p} minimum aberration design consisting of the design factors only. This provides an easy way to construct sliced minimum aberration designs by using existing minimum aberration designs of fewer number of factors. To construct 2^{k+1-p} sliced minimum aberration designs, one can begin by generating a

2^{k-p} minimum aberration design consisting of the design factors for one platform and then repeat the same design for the other platform. To illustrate this point, we now use Theorem 1 to construct a sliced minimum aberration design for Example 1.

EXAMPLE 1 (CONTINUED). The design $\mathbf{I} = \mathbf{123}$ is a 2^{3-1} minimum aberration design consisting of design factors that can be used for both platforms. Multiplying $\mathbf{I} = \mathbf{123}$ by S provides the sliced defining relation $\mathbf{S} = \mathbf{123S}$ of the sliced minimum aberration design $d^{(2)}$. As $d^{(2)}$ has the defining relation $\mathbf{I} = \mathbf{123}$, it is not a 2^{4-1} minimum aberration design consisting of all four factors. Design $d^{(2)}$ has one sliced word $\mathbf{123S}$ of length four. The sliced wordlength pattern is then $SW(d^{(2)}) = (4^1)$. To illustrate that $d^{(2)}$ is a better design than $d^{(1)}$, we compare their confounded effects involving the slice factor S . These confounded effects for $d^{(1)}$ are $\mathbf{S} = \mathbf{123}$, $\mathbf{1S} = \mathbf{23}$, $\mathbf{2S} = \mathbf{13}$, and $\mathbf{3S} = \mathbf{12}$. In contrast, the confounded effects of $d^{(2)}$ are $\mathbf{S} = \mathbf{123S}$, $\mathbf{1S} = \mathbf{23S}$, $\mathbf{2S} = \mathbf{13S}$, and $\mathbf{3S} = \mathbf{12S}$. The design $d^{(2)}$ provides more estimation capacity for the slice factor S because, according to sliced effect hierarchy principle, $\mathbf{123S}$ is less likely to be important than $\mathbf{123}$. Along the same lines, $d^{(2)}$ confounds $\mathbf{1S}$ with $\mathbf{23S}$ and this provides more estimation capacity than confounding $\mathbf{1S}$ with $\mathbf{23}$. Using the same argument, it is more desirable to confound $\mathbf{2S}$ with $\mathbf{13S}$ (vs. $\mathbf{13}$ in $d^{(1)}$) and $\mathbf{3S}$ with $\mathbf{12S}$ (vs. $\mathbf{12}$ in $d^{(1)}$).

To construct a sliced minimum aberration design for the complete design d in Definition 1, we use the idea in Theorem 1 to provide a simple algorithm as follows:

ALGORITHM 1.

- Step 1: For platform P_1 of the experiment in Definition 1, find a minimum aberration design for the design factors and use it for all sub designs d_1, d_2, \dots, d_s .*
- Step 2: For $j = 1, \dots, s$, add the j th level of the slice factor to experimental versions of d_j .*
- Step 3: Combine all sub designs in Step 2 to obtain a sliced minimum aberration design d .*

Next we use an example involving a greater number of factors to illustrate Algorithm 1.

EXAMPLE 2. For the experiment in Definition 1, let $k = 8$ and $s = 2$ for two platform P_1 and P_2 . The slice factor S is defined with two levels $+$ and $-$, representing platforms P_1 and P_2 , respectively. Consider a 2^{9-3} complete design d that consists of d_1 and d_2 , each being a 2^{8-3} factorial design. Following Step 1 of Algorithm 1, three generators g_1 : **13458**, g_2 : **1247**, and g_3 : **1236** provide the 2^{8-3} minimum aberration design of the design factors (Wu and Hamada 2011, p. 254). Using the same design for both platforms results in the defining relation of $\mathbf{I} = \mathbf{1236} = \mathbf{1247} = \mathbf{3467} = \mathbf{13458} = \mathbf{24568} = \mathbf{23578} = \mathbf{15678}$ for both d_1 and d_2 .

The words **3467**, **24568**, **23578**, and **15678** are simply obtained by the products $g_2 \times g_3$, $g_1 \times g_3$, $g_1 \times g_2$, and $g_1 \times g_2 \times g_3$ respectively. The resulting experimental versions for both d_1 and d_2 are $\{8, 167, 2678, 12, 36, 1378, 237, 12368, 47, 1468, 246, 12478, 34678, 134, 2348, 123467, 5, 15678, 2567, 1258, 3568, 1357, 23578, 12356, 4578, 1456, 24568, 12457, 34567, 13458, 2345, 12345678\}$, where 24568, for instance, represents the version that has five factors 2, 4, 5, 6, and 8 at $+$ and the other three factors 1, 3, and 7 at $-$. Adding S to all experimental versions of d_2 and combining it with d_1 provides the 2^{9-3} sliced minimum aberration design with the sliced defining relation $\mathbf{S} = \mathbf{1236S} = \mathbf{1247S} = \mathbf{3467S} = \mathbf{13458S} = \mathbf{24568S} = \mathbf{23578S} = \mathbf{15678S}$ that has the sliced wordlength pattern $(5^3, 6^4)$.

Algorithm 1 is for the case where all the combinations of design factors are feasible for all the platforms. Next, we propose a procedure for the more general situation where some combinations of the design factors may be infeasible for some platforms. This incorporates a typical design constraint wherein not all factors combinations may be feasible for all platforms. An example of such a design constraint on a smartwatch would be the inability to include two image based factors and a brand logo on one screen. The same combination may be feasible on a desktop. For situations where some combinations of the design factors may be infeasible for some platforms, exact sliced minimum aberration designs cannot be guaranteed. We revisit Example 2 to illustrate this.

EXAMPLE 2 (CONTINUED). Let us assume that version 8 should be in sub design d_1 and the combination 24568 cannot be used in the experimental versions of sub design d_2 . Because version

8 is required on platform P_1 , the 2^{8-3} minimum aberration design provided by g_1 : **8 = 1345**, g_2 : **7 = 124**, and g_3 : **6 = 123** is used for sub design d_1 . However, the same design cannot be used for sub design d_2 as the experimental versions 24568 and 12345678 include the combination 24568. Therefore, the sliced minimum aberration design cannot be obtained.

For the situation where some combinations of the design factors may be infeasible for some platforms, sliced aberration based criterion can be used to search all designs that are orthogonal and balanced in main effects until it reaches the best possible sliced factorial design. For ease of explanation, let us assume $s = 2$. By applying Step 1 of Algorithm 1, there are 2^p fractions to be selected for sub design d_1 . It is assumed that all experimental versions of at least one of the fractions are feasible within each platform. We first introduce some useful notation.

For the experiment in Definition 1 with $s = 2$, let M be the set of all 2^{k+1-p} complete designs d 's. We partition the set M into three disjoint sets, M_1 , M_2 , and M_3 , based on the defining relations of the sub designs d_j 's:

M_1 : The set of 2^{k+1-p} complete designs d 's for which the defining relations of both sub designs d_j 's are the same.

M_2 : The set of 2^{k+1-p} complete designs d 's for which the defining relations of both sub designs d_j 's have the same words but different use of $+/-$.

M_3 : The set of 2^{k+1-p} complete designs d 's for which the defining relations of sub designs d_j 's have different words.

A sliced minimum aberration design generated by Algorithm 1 belongs to the set M_1 . Any design from the set M_3 is neither balanced nor orthogonal in main effects. Therefore, no designs from the set M_3 is considered. For the case where some combinations of design factors may be infeasible and no sliced minimum aberration design from the set M_1 is possible, one should search for a design from set M_2 . We next illustrate how to construct a design from set M_2 by revisiting Example 2.

EXAMPLE 2 (CONTINUED). Construct d_1 using three generators g_1 : **13458**, g_2 : **1247**, and g_3 : **1236**, and d_2 using three generators $-g_1$: **-13458**, g_2 : **1247**, and g_3 : **1236**. As a result, d_1 and

d_2 have defining relations $\mathbf{I} = \mathbf{1236} = \mathbf{1247} = \mathbf{3467} = \mathbf{13458} = \mathbf{24568} = \mathbf{23578} = \mathbf{15678}$ and $\mathbf{I} = \mathbf{1236} = \mathbf{1247} = \mathbf{3467} = -\mathbf{13458} = -\mathbf{24568} = -\mathbf{23578} = -\mathbf{15678}$, respectively. By combining d_1 and d_2 , the sliced defining relation $\mathbf{S} = \mathbf{1236S} = \mathbf{1247S} = \mathbf{3467S} = \mathbf{13458} = \mathbf{24568} = \mathbf{23578} = \mathbf{15678}$ is provided for d whose slicing is determined by g_1 and has the sliced wordlength pattern (5^7) - compare with the sliced wordlength pattern $(5^3, 6^4)$ of the sliced minimum aberration design. This design belongs to set M_2 as the defining relations of both its sub designs have the same words but with different use of $+/-$. Similarly, the design whose slicing is determined by g_1, g_2 is constructed by using three generators g_1, g_2 , and g_3 for d_1 and three generators $-g_1, -g_2$, and g_3 for d_2 . The design whose slicing is determined by g_1, g_2 belongs to set M_2 .

The sliced aberration based criterion can be used to rank all the designs in set M_2 which allows a practitioner to search for the best possible sliced factorial design when the exact sliced minimum aberration designs is not guaranteed. To illustrate this procedure, we revisit Example 2.

EXAMPLE 2 (CONTINUED). As we discussed earlier, the sliced minimum aberration design cannot be obtained for the case where version 8 is required in sub design d_1 and the combination 24568 cannot be used in the experimental versions of sub design d_2 . The design whose slicing is determined by g_1 provides the experimental versions {NULL, 1678, 267, 128, 368, 137, 2378, 1236, 478, 146, 2468, 1247, 3467, 1348, 234, 1234678, 58, 1567, 25678, 125, 356, 13578, 2357, 123568, 457, 14568, 2456, 124578, 345678, 1345, 23458, 1234567} for d_2 , where NULL represents the version with all eight design factors at $-$. The design whose slicing is determined by g_3 can also be used that provides the experimental versions {68, 17, 278, 126, 3, 13678, 2367, 1238, 467, 148, 24, 124678, 3478, 1346, 23468, 12347, 56, 1578, 257, 12568, 358, 13567, 235678, 1235, 45678, 145, 2458, 124567, 3457, 134568, 23456, 1234578} for d_2 . The design whose slicing is determined by g_2 results in the experimental versions {87, 16, 268, 127, 367, 138, 23, 123678, 4, 14678, 2467, 1248, 3468, 1347, 23478, 12346, 57, 1568, 256, 12578, 35678, 135, 2358, 123567, 458, 14567, 245678, 1245, 3456, 134578, 23457, 1234568} for d_2 . Therefore, the design whose slicing is determined by g_2 cannot be used as the experimental version 245678 includes the combination 24568.

Comparing the sliced wordlength pattern (5^7) of the design whose slicing is determined by g_1 with the sliced wordlength pattern $(4^2, 5^3, 6^2)$ of the design whose slicing is determined by g_3 , shows that the design whose slicing is determined by g_1 is preferred since its shortest sliced word is of length 5 which is greater than 4.

Table 4 includes all different designs whose slicing is determined by different combinations of three generators g_1 , g_2 , and g_3 . All eight sliced factorial designs considered in Table 4 are the designs that best satisfy Property 1 because all their sub designs are minimum aberration designs consisting of design factors only. However, these sliced factorial designs perform differently in terms of Property 2. Among all these sliced factorial designs that are orthogonal and balanced in main effects, the sliced minimum aberration design is the best.

Table 4 Rank of 2^{9-3} Orthogonal Sliced Factorial Designs for Example 2 with Three Generators g_1 : 13458, g_2 : 1247, and g_3 : 1236

Rank	Sliced Factorial Design	< determining generators >	Sliced Wordlength Pattern	Comment
1	Sliced minimum aberration design	Empty set	$(5^3, 6^4)$	
2	Slicing determined by g_1	$\langle g_1 \rangle = \{13458, 24568, 23578, 15678\}$	(5^7)	
3	Slicing determined by g_2	$\langle g_2 \rangle = \{1247, 3467, 23578, 15678\}$	$(4^2, 5^3, 6^2)$	MA* design
3	Slicing determined by g_1, g_2	$\langle g_1, g_2 \rangle = \{1247, 3467, 13458, 24568\}$	$(4^2, 5^3, 6^2)$	MA design
3	Slicing determined by g_3	$\langle g_3 \rangle = \{1236, 3467, 24568, 15678\}$	$(4^2, 5^3, 6^2)$	MA design
3	Slicing determined by g_1, g_3	$\langle g_1, g_3 \rangle = \{1236, 3467, 13458, 23578\}$	$(4^2, 5^3, 6^2)$	MA design
3	Slicing determined by g_2, g_3	$\langle g_2, g_3 \rangle = \{1236, 1247, 24568, 23578\}$	$(4^2, 5^3, 6^2)$	MA design
3	Slicing determined by g_1, g_2, g_3	$\langle g_1, g_2, g_3 \rangle = \{1236, 1247, 13458, 15678\}$	$(4^2, 5^3, 6^2)$	MA design

* MA: minimum aberration

5. Empirical Application

The purpose of this study was to improve the email design for a digital magazine published by the business school of a university. This magazine has both print and digital versions. The business school sends an email to its alumni base twice a year inviting them to read a variety of news in the magazine. The invitation email contains links to the landing page and multiple articles that relate to the business school. Version one in Table 6 shows the top half of the most recent issue of the magazine prior to this study. To maintain integrity of the review process we have masked parts of the magazine that may reveal the school name. We intend to remove the masked sections at a later time.

For this study, the team managing the digital magazine identified six binary design factors for the multivariate test for two platforms P_1 and P_2 . Platform P_1 refers to a tablet or cell phone and P_2 refers to a desktop or laptop computer. We decided to combine the two devices for each platform in order to have enough observations/platform. For contexts where the expected sample size is large, this aggregation would be unnecessary. The slice factor S is defined as a binary for which the level $-$ represents platform P_1 and the level $+$ represents platform P_2 . For a full design we would have to create 2^6 versions for each platform P_1 and P_2 . Instead, using the design criteria developed in Section 4, we create 2^3 versions to perform the multivariate testing. The design we use is a 2^{7-3} sliced factorial design. Table 5 includes six binary design factors that were identified for this study. These factors are 1: full width banner, 2: stories with minimal teaser, 3: class note stories, 4: cover image at top, 5: full width logo, and 6: call to action button. For each design factor, the $+$ level is the change from the control version of magazine. The control version was identical to the most recently published magazine prior to this study (version one in Table 6).

Each additional version in a multivariate design has an associated incremental cost. In our empirical application, because of the extensive programming related resource constraints, the maximum number of versions was restricted to eight. In another multi-platform experiment context, such as website redesign, constraints include number of websites that one could realistically test

Table 5 Six Binary Design Factors for Update Email Optimization

Factor	+	–
1	<p>Selected articles in a full width banner format</p> 	<p>Keep as is in the control version</p> 
2	<p>Stories with very minimal teaser copy</p> 	<p>Keep as is in the control version</p> 
3	<p>The class notes story as a set of 3 example notes</p> <p>Class Notes</p> <p>Business from around the world are earning promotions, accepting new positions, starting new businesses, and marking major life events. Read the latest alumni news in Class Notes.</p> <div> <p>with multinational corporations on complicated CFO and CIO issues such as cost reduction and operations improvement ... Read more »</p> <p>director for Sanrio, a former Target buyer, current member of the Business Alumni Board, and a contemporary ... Read more »</p> <p>founded I with her husband I to inspire two-wheeled tourism in North America's best biking cities. The website ... Read more »</p> </div>	<p>Keep as is in the control version</p> 
4	<p>An image of the cover at top</p> <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p> 	<p>Keep as is in the control version</p>
5	<p>The full width Logo on top</p> 	<p>Keep as is in the control version</p> 
6	<p>A call to action button on top</p> <p>Read the Spring 2015 Issue »</p>	<p>Keep as is in the control version</p>

without disrupting business. The eight versions in our setup included the control version, the format used in the most recent magazine prior to this study. We used Algorithm 1 to generate a sliced factorial design for this study. Following Step 1 of Algorithm 1, three generators **124**, **135**, and **236** provide the 2^{6-3} minimum aberration design of the design factors (Wu and Hamada 2011, p. 253). Using the same design for both platforms results in the defining relation of $\mathbf{I} = \mathbf{124} = \mathbf{135} = \mathbf{236} = \mathbf{456} = \mathbf{2345} = \mathbf{1346} = \mathbf{1256}$ for both d_1 and d_2 . Both sub designs d_1 and d_2 are 2^{6-3} minimum aberration designs with resolution III. The resulting experimental versions for both d_1 and d_2 are $\{NULL, 145, 246, 1256, 356, 1346, 2345, 123\}$, where *NULL* is the control version and 145, for instance, represents the version that has three factors 1, 4, and 5 at + and the other three factors 2, 3, and 6 at -. Adding S to all experimental versions of d_2 and combining it with d_1 provides the 2^{7-3} sliced minimum aberration design with the sliced defining relation $\mathbf{S} = \mathbf{124S} = \mathbf{135S} = \mathbf{236S} = \mathbf{456S} = \mathbf{2345S} = \mathbf{1346S} = \mathbf{1256S}$ that has the sliced wordlength pattern $(4^4, 5^3)$. Algorithm 1 generated a sliced minimum aberration design that required eight versions. In contrast, a blocked design would have required sixteen versions thus needing greater resources compared to our methodology.

The first version of our design is the control version which is partially presented in Table 6. Version two has factors 1, 4, and 5 that are at + levels and the remaining three factors are at - levels. Similarly, version three has factors 2, 4, and 6 at + levels although the other three factors are at - levels. Version four has factors 1, 2, 5, and 6 at + levels and the remaining factors are at - levels. Version five has factors 3, 5, and 6 at + levels and the rest at - levels. Version six has factors 1, 3, 4, and 6 at + levels and the rest at - levels. Version seven has factors 2, 3, 4, and 5 at + levels and the rest at - levels. Finally for version eight, three factors 1, 2, and 3 are at + levels and the remaining factors are at - levels. Table 6 and Table 7 include the top part of all eight versions used for our empirical application. For versions five, six, seven, and eight class notes are displayed as a set of 3 example notes at the bottom.

The response variable for this study is the number of page views that are obtained from

Table 6 Versions One, Two, Three, and Four - Update Spring 2015







Version one: control version	Version two: 145
1 : -, 2 : -, 3 : -, 4 : -, 5 : -, 6 : -	1 : +, 2 : -, 3 : -, 4 : +, 5 : +, 6 : -
 <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p> <p>Alumni Inspiration: 8 to Watch The impact of alumni is far-reaching. Meet eight Business whose influence will be felt in their industries and beyond for years to come. Read about these amazing alumni.</p> <p>What Inspires You? and I share insight into why teamwork matters. Find out what inspires the Poytingers.</p> <p>Research: Now You Know When is a good time to start talking to children about money? How does the long-term value of American-made Toyotas compare to those made in Japan? What makes for a kinder, gentler workplace? Get answers to these intriguing questions.</p>	 <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p>  <p>» 8 to Watch</p> <p>Alumni Inspiration: 8 to Watch The impact of alumni is far-reaching. Meet eight Business whose influence will be felt in their industries and beyond for years to come. Read about these amazing alumni.</p> <p>What Inspires You? teamwork matters. Find out what inspires the Poytingers.</p>
Version three: 246	Version four: 1256
1 : -, 2 : +, 3 : -, 4 : +, 5 : -, 6 : +	1 : +, 2 : +, 3 : -, 4 : -, 5 : +, 6 : +
 <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p>  <p>Read the Spring 2015 Issue »</p> <p>Alumni Inspiration: 8 to Watch Meet eight alumni who are making an impact.</p> <p>What Inspires You? Find out what inspires the co-owners of Splash Studio in</p> <p>Research: Now You Know Get answers that you can trust to important business questions.</p>	 <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p> <p>Read the Spring 2015 Issue »</p> <p>» 8 to Watch</p> <p>Alumni Inspiration: 8 to Watch Meet eight alumni who are making an impact.</p> <p>What Inspires You? Find out what inspires the co-owners of Splash Studio in</p> <p>Research: Now You Know Get answers that you can trust to important business questions.</p>

Table 7 Versions Five, Six, Seven, and Eight - Update Spring 2015

Version five: 356	Version six: 1346
1 : -, 2 : -, 3 : +, 4 : -, 5 : +, 6 : +	1 : +, 2 : -, 3 : +, 4 : +, 5 : -, 6 : +
<p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p> <p>Read the Spring 2015 Issue »</p>  <p>Alumni Inspiration: 8 to Watch</p> <p>The impact of alumni is far-reaching. Meet eight Business whose influence will be felt in their industries and beyond for years to come. Read about these amazing alumni.</p>	 <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p>  <p>Read the Spring 2015 Issue »</p>  <p>Alumni Inspiration: 8 to Watch</p> <p>The impact of alumni is far-reaching. Meet eight Business whose influence will be felt in their industries and beyond for years to come. Read about these amazing alumni.</p>
Version seven: 2345	Version eight: 123
1 : -, 2 : +, 3 : +, 4 : +, 5 : +, 6 : -	1 : +, 2 : +, 3 : +, 4 : -, 5 : -, 6 : -
<p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p>  <p>Alumni Inspiration: 8 to Watch</p> <p>Meet eight alumni who are making an impact.</p> 	 <p>See the Latest Alumni News in the Spring 2015 Issue of</p> <p>Find out what's happening with your fellow alumni, and read the inspiring stories coming out of the School of Business today.</p>  <p>Alumni Inspiration: 8 to Watch</p> <p>Meet eight alumni who are making an impact.</p>  <p>What Inspires You?</p> <p>Find out what inspires the co-owners of Splash Studio in</p>

Google Analytics. Google Analytics records page views for each version of the study. These data are aggregated across users that are exposed to each version. How page views vary within a version is not known to us. This is very typical of Google Analytics data. Therefore, a popular method called Lenth’s test (Lenth 1989) is used to identify statistically significant factors for this study. The Lenth’s test is specifically designed for testing effects in experiments for which variance estimates are not available. It is simple to use and performs well according to Hamada and Balakrishnan (1998) who report an extensive review of different methods one could use to analyze an unreplicated factorial experiment. The Lenth’s test is also well studied by Ye and Hamada (2000).

Table 8 includes some descriptive statistics for this study. The total number of recipients is 25,693 and they are divided into roughly equal eight sets, each set of recipients received a version of the email.

Table 8 Descriptive Statistics for Empirical Application

	V1	V2	V3	V4	V5	V6	V7	V8
Recipients	3225	3224	3205	3205	3211	3215	3209	3199
page views	674	327	402	350	587	304	580	337

There are eight versions for each platform and as a result seven effects of design factors can be estimated per platform. Table 9 includes the confounded effects within each platform of our empirical study. These confounded effects are identified by the sliced factorial design we generated by using Algorithm 1. Such confounds are implicit in any fractional factorial design that one uses in marketing and we make them explicit to offer greater clarity. For convenience, we label each set of confounded effects.

In the sliced factorial design framework, slices are used for analyses on both P_1 and P_2 . Tables 10 and 11 include the effects of design factors that are estimated using sub designs d_1 on platform P_1 and d_2 on platform P_2 , respectively. Lenth’s method is used to test the significance of effects and to report the p-values. Following Property 1, the analyses for each slice are used to estimate the effect of designs factors within each platform.

Table 9 Confounded Effects

Labels	Confounded Effects
A	1 = 24 = 35 = 346 = 256 = 1236 = 1456 = 12345
B	2 = 14 = 36 = 345 = 156 = 2456 = 1235 = 12346
C	3 = 15 = 26 = 245 = 146 = 1234 = 3456 = 12356
D	4 = 12 = 56 = 235 = 136 = 1345 = 2346 = 12456
E	5 = 13 = 46 = 126 = 234 = 1245 = 2356 = 13456
F	6 = 23 = 45 = 134 = 125 = 1246 = 1356 = 23456
G	16 = 34 = 25 = 145 = 246 = 356 = 123 = 123456

Table 10 Results for P_1

Effect	Estimate	P-value
A	0.010	> 0.2
B	0.006	> 0.2
C	0.011	> 0.2
D	0.001	> 0.2
E	-0.025	0.12
F	0.038	0.04 *
G	-0.044	0.03 *

Table 11 Results for P_2

Effect	Estimate	P-value
A	0.181	> 0.2
B	0.182	> 0.2
C	0.278	0.08 *
D	0.062	> 0.2
E	-0.068	> 0.2
F	-0.055	> 0.2
G	-0.102	> 0.2

Comparing the results of Tables 10 and 11 suggests that effect **C** is significant on P_2 (desktops and laptops) although two effects **F** and **G** are significant on P_1 (smartphone, tablet). Table 9 reveals that effect **C** is the sum of the following confounded effects **3, 15, 26, 245, 146, 1234, 3456, 12356**. As slices follow effect hierarchy principle, **C** can be viewed to represent effect **3** by assuming that all higher-order confounded effects are negligible. The main takeaway for P_2 from Table 11 is that displaying class notes as a set of three (see factor 3 in Table 5) will improve the number of page views and that other factors are unlikely to have a positive effects on page views.

Looking at the results for P_1 in Tables 10, the two significant effects are **F** and **G**. Table 9 shows that effect **F** is the sum of the following confounded effects **6, 23, 45, 134, 125, 1246, 1356, 23456**. As slices follow effect hierarchy principle, **F** can be viewed to represent effect **6** by assuming that all higher-order confounded effects are negligible. Further, effect **G** is the sum of the following confounded effects **16, 34, 25, 145, 246, 356, 123, 123456**. Once again, the effect hierarchy principle suggests that **G** represents sum of three effects **16, 34, and 25**. Next, by applying the effect

heredity principle, **G** can be viewed to represent effect **16** as its parent **6** is significant. The effects **34** and **25** can be safely assumed to be zero because their parent effects (**2**, **3**, **4**, **5**) are statistically insignificant. The effect **6** is positive and the effect **16** is negative. Therefore, the main takeaway for P_1 is that adding the call to action button on top (factor 6 in Table 5) and refraining from the full width banner format (see factor 1 in Table 5) will improve number of page views. The remaining factors do not impact page views in any way so should not be considered for adoption the next digital magazine.

Following Property 2, to compare the results of two platforms, the complete design d is then used to estimate the slice factor and its interaction with the platform specific significant effects. Table 12 includes four effects **S**, **3S**, **6S**, and **16S**. The effect **S** is positive and significant implying that P_2 brings significantly more page views compared to P_1 . Further, the magnitude of the effect **S** is around three times larger than the effects of design factors. This is important to note because it is consistent with the sliced effect hierarchy principle. The effects **3S**, **6S**, and **16S** uncover the way the effects **3**, **6**, and **16** interact with the slice factor, respectively, meaning how these effects differentially affects page views from P_1 to P_2 . We note that these interactions are not statistically significant implying that the differential effects of **3**, **6**, and **16** on page views from P_1 to P_2 are not significant. As a result, one version that adopts the changes of displaying “class notes as a set of three”, adding the “call to action button”, and refraining from the “full width banner format” will improve number of page views for both platforms P_1 and P_2 .

Table 12 **Slice factor Behavior**

Effect	Estimate	P-value
S	0.671	< 0.01 *
3S	0.134	0.14
6S	−0.046	> 0.2
16S	−0.029	> 0.2

Next, we calculate the expected incremental gain in page views because of the design changes the multi-platform multivariate test reveals. For P_2 , average pageview can be estimated using

$$\text{Average Pageview} = 1.864 + 0.139\mathbf{C}. \quad (7)$$

When compared to the control group, this suggests an expected 16% gain in page views when the “class notes story” ($\mathbf{C} = +$ vs. $\mathbf{C} = -$) is included in the magazine design. Similarly for P_1 , average pageview can be estimated by using

$$\text{Average Pageview} = 1.193 + 0.019\mathbf{F} - 0.022\mathbf{A}\mathbf{F}. \quad (8)$$

When compared to the control group, this provides a 7% gain in page views when the “call to action” button is added on the top and the “full width banner format” ($\mathbf{F} = +$ and $\mathbf{A} = -$ vs. $\mathbf{F} = -$ and $\mathbf{A} = -$) is excluded in the magazine design.

In sum, the sliced factorial design methodology suggests different set of designs factors would increase page views for P_2 and P_1 . Displaying the class notes story will improve pageview on P_2 and adding the call to action button on top while avoiding the full width banner format will improve page views on P_1 . As a result of this multivariate test, the next issue of the digital magazine (see Figure 3) adopted the following changes to increase page views: (i) display the class notes story, (ii) add the call to action button on top, and (iii) not use the full width banner format.

6. Conclusion

A unique aspect of multivariate testing in the online space is that testing needs to be conducted across platforms that include a desktop, a laptop, a tablet, and a smartphone. The existing design literature does not offer precise guidance for such a multi-platform context. The primary focus of this paper is to fill this void in the literature. Our primary contribution is that we develop a multi-platform, multivariate design that allows us to uncover effects for each platform and compare test results across different platforms by studying the interactions of design factors and platforms. We develop a new factorial design, called the sliced factorial design, that is used to perform a multivariate test within each platform.

We propose a novel sliced effect hierarchy principle that generalizes the widely used effect hierarchy principle to the multi-platform context. We show that widely used resolution and aberration

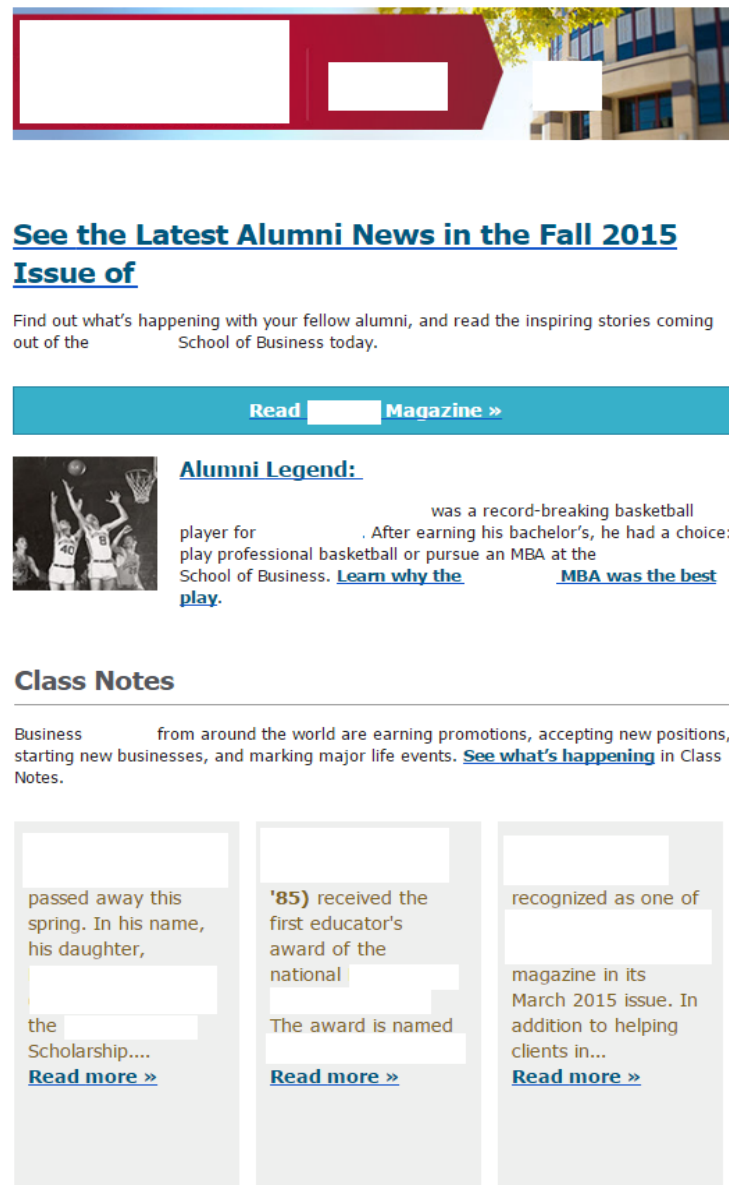


Figure 3 Next Issue of the Magazine After the Study- Fall 2015.

criteria fail to satisfy the sliced effect hierarchy principle. To address this problem, we propose sliced resolution and sliced aberration as the design criteria for our multi-platform context. It is well known that minimum aberration designs are popular among practitioners and tables to construct them are readily available in software and textbooks. In an effort to build upon the rich results of the aberration literature, we prove a theorem that connects sliced factorial designs to minimum aberration designs. This theorem helps construct sliced factorial designs that we propose. Using

the novel design criteria and this theorem, we develop two algorithms to construct sliced factorial designs. The first algorithm is for the case where all combinations of design factors are feasible for all platforms. This algorithm works for both symmetric and asymmetric designs. The second algorithm is usable in situations where some combinations of design factors may be infeasible for some platforms.

We find that it is desirable to have the sliced factorial design be divided into homogeneous slices: experimental versions made for a slice should be as similar as possible to the ones for other slices. In terms of resources, the multivariate design we propose requires a smaller number of versions compared to other designs. From the standpoint of cost, this is highly desirable aspect of the designs we propose. We illustrate that although slices are used to uncover factor effects within each platform, the sliced factorial design can be used to compare the results across different platforms.

Finally, we illustrate our novel design framework in the context of an empirical email optimization application intended to improve engagement for a digital magazine. Algorithm 1 was used to generate a sliced factorial design for this study. Google Analytics was used to record the number of page views for each version. Lenth's test, which is well suited for such data, revealed interesting insight about how different factors were effective for different platforms. Based on the results for this multi-platform experiment, the expected gain in page views for the two platforms was 16% and 7%.

In closing, we hope that the proposed design framework will become a useful way for practitioners to design and analyze multi-platform experiments online.

References

- Ba, S., Myers, W. R., and Breneman, W. A. (2015). Optimal sliced latin hypercube designs. *Technometrics*, 57(4):479–487.
- Box, G. E., Hunter, W. G., Hunter, J. S., et al. (1978). *Statistics for experimenters*. John Wiley & Sons.
- Chaffey, D. (2016). Mobile marketing statistics compilation, available at <http://www.smartinsights.com/mobile-marketing/mobile-marketing-analytics/mobile-marketing-statistics/>.

- Cheng, C.-S., Steinberg, D. M., and Sun, D. X. (1999). Minimum aberration and model robustness for two-level fractional factorial designs. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 61(1):85–93.
- Hamada, M. and Balakrishnan, N. (1998). Analyzing unreplicated factorial experiments: a review with some new proposals. *Statistica Sinica*, pages 1–28.
- Huang, H., Yang, J.-F., and Liu, M.-Q. (2014). Construction of sliced (nearly) orthogonal latin hypercube designs. *Journal of Complexity*, 30(3):355–365.
- Lenth, R. V. (1989). Quick and easy analysis of unreplicated factorials. *Technometrics*, 31(4):469–473.
- Qian, P. Z. and Wu, C. J. (2009). Sliced space-filling designs. *Biometrika*, pages 945–956.
- Qian, P. Z. G. (2012). Sliced latin hypercube designs. *Journal of the American Statistical Association*, 107(497):393–399.
- Stanhope, J. (2013). The forrester wave: Online testing platforms, q1 2013. Technical report, Forrester Research.
- Wu, C. J. and Hamada, M. S. (2011). *Experiments: planning, analysis, and optimization*, volume 552. John Wiley & Sons.
- Xu, X., Haaland, B., and Qian, P. Z. G. (2011). Sudoku-based space-filling designs. *Biometrika*, 98(3):711–720.
- Yang, J.-F., Lin, C. D., Qian, P. Z., and Lin, D. K. (2013). Construction of sliced orthogonal latin hypercube designs. *Statistica Sinica*, pages 1117–1130.
- Ye, K. Q. and Hamada, M. (2000). Critical values of the lenth method for unreplicated factorial designs. *Journal of Quality Technology*, 32(1):57–66.
- Yin, Y., Lin, D. K., and Liu, M.-Q. (2014). Sliced latin hypercube designs via orthogonal arrays. *Journal of Statistical Planning and Inference*, 149:162–171.